

Dec 11, 2025

GPT-5.2 Prompting Guide

This recipe is archived and may reference outdated models or APIs.



Mandeep Singh (OpenAI) , Emre Okcular (OpenAI)

[View on GitHub](#)[Download raw](#)

1. Introduction

GPT-5.2 is our newest flagship model for enterprise and agentic workloads, designed to deliver higher accuracy, stronger instruction following, and more disciplined execution across complex workflows. Building on GPT-5.1, GPT-5.2 improves token efficiency on medium-to-complex tasks, produces cleaner formatting with less unnecessary verbosity, and shows clear gains in structured reasoning, tool grounding, and multimodal understanding.

GPT-5.2 is especially well-suited for production agents that prioritize reliability, evaluability, and consistent behavior. It performs strongly across coding, document analysis, finance, and multi-tool agentic scenarios, often matching or exceeding leading models on task completion. At the same time, it remains prompt-sensitive and highly steerable in tone, verbosity, and output shape, making explicit prompting an important part of successful deployments.

While GPT-5.2 works well out of the box for many use cases, this guide focuses on prompt patterns and migration practices that maximize performance in real production systems. These recommendations are drawn from internal testing and customer feedback, where small changes to prompt structure, verbosity constraints, and reasoning settings often translate into large gains in correctness, latency, and developer trust.

2. Key behavioral differences

[Ask AI](#)

Loading docs agent...

Example clamp adjust based on desired verbosity:

```
<output_verbosity_spec>
```

- Default: 3–6 sentences or ≤ 5 bullets for typical answers.
- For simple “yes/no + short explanation” questions: ≤ 2 sentences.
- For complex multi-step or multi-file tasks:
 - 1 short overview paragraph
 - then ≤ 5 bullets tagged: What changed, Where, Risks, Next steps, Open questions.
- Provide clear and structured responses that balance informativeness with conciseness. Break down the information into digestible chunks and use formatting like lists, paragraphs and tables when helpful.
- Avoid long narrative paragraphs; prefer compact bullets and short sections.



Loading docs agent...

dense contexts.

<long_context_handling>



- For inputs longer than ~10k tokens (multi-chapter docs, long threads, multiple PDFs):
 - First, produce a short internal outline of the key sections relevant to the user's request.
 - Re-state the user's constraints explicitly (e.g., jurisdiction, date range, product, team) before answering.
 - In your answer, anchor claims to sections ("In the 'Data Retention' section...") rather than speaking generically.
 - If the answer depends on fine details (dates, thresholds, clauses), quote or paraphrase them.
- </long_context_handling>

3.4 Handling ambiguity & hallucination risk

Loading docs agent...

- Specific numbers or claims not grounded in context,
 - Overly strong language (“always,” “guaranteed,” etc.).
 - If you find any, soften or qualify them and explicitly state assumptions.
- </high_risk_self_check>

4. Compaction (Extending Effective Context)

For long-running, tool-heavy workflows that exceed the standard context window, GPT-5.2 with Reasoning supports response compaction via the `/responses/compact` endpoint. Compaction performs a loss-aware compression pass over prior conversation state, returning encrypted, opaque items that preserve task-relevant information while dramatically reducing token footprint. This allows the model to continue reasoning across extended workflows without hitting context limits.

Loading docs agent...

- Monitor context usage and plan ahead to avoid hitting context window limits
- Compact after major milestones (e.g., tool-heavy phases), not every turn
- Keep prompts functionally identical when resuming to avoid behavior drift
- Treat compacted items as opaque; don't parse or depend on internals

For guidance on when and how to compact in production, see the [Conversation State](#) guide and [Compact a Response](#) page.

Here is an example:

```
1 from openai import OpenAI
2 import json
3
4
5 client = OpenAI()
6
```



Loading docs agent...

```
indent=2))
```

5. Agentic steerability & user updates

GPT-5.2 is strong on agentic scaffolding and multi-step execution when prompted well. You can reuse your GPT-5.1 `<user_updates_spec>` and `<solution_persistence>` blocks.

Two key tweaks could be added to further push the performance of GPT-5.2:

- Clamp verbosity of updates (shorter, more focused).
- Make scope discipline explicit (don't expand problem surface area).

Example updated spec:

Loading docs agent...

configs, logs).

- You reference specific IDs, URLs, or document titles.
- Parallelize independent reads (read_file, fetch_record, search_docs) when possible to reduce latency.
- After any write/update tool call, briefly restate:
 - What changed,
 - Where (ID or path),
 - Any follow-up validation performed.

</tool_usage_rules>

7. Structured extraction, PDF, and Office workflows

This is an area where GPT-5.2 clearly shows strong improvements. To get the most out of it:

Loading docs agent...

8. Prompt Migration Guide to GPT 5.2

This section helps you migrate prompts and model configs to GPT-5.2 while keeping behavior stable and cost/latency predictable. GPT-5-class models support a `reasoning_effort` knob (e.g., `none`|`minimal`|`low`|`medium`|`high`|`xhigh`) that trades off speed/cost vs. deeper reasoning.

Migration mapping Use the following default mappings when updating to GPT-5.2

Current model	Target model	Target reasoning_effort	Notes
GPT-4o	GPT-5.2	none	Treat 4o/4.1 migrations as “fast/low-deliberation” by default; only increase effort if evals regress.

Loading docs agent...

- Step 4: If regressions, tune the prompt. Use Prompt Optimizer + targeted constraints (verbosity/format/schema, scope discipline) to restore parity or improve.
- Step 5: Re-run Evals after each small change. Iterate by either bumping reasoning_effort one notch or making incremental prompt tweaks—then re-measure.

9. Web search and research

GPT-5.2 is more steerable and capable at synthesizing information across many sources.

Best practices to follow:

Loading docs agent...

10. Conclusion

GPT-5.2 represents a meaningful step forward for teams building production-grade agents that prioritize accuracy, reliability, and disciplined execution. It delivers stronger instruction following, cleaner output, and more consistent behavior across complex, tool-heavy workflows. Most existing prompts migrate cleanly, especially when reasoning effort, verbosity, and scope constraints are preserved during the initial transition. Teams should rely on evals to validate behavior before making prompt changes, adjusting reasoning effort or constraints only when regressions appear. With explicit prompting and measured iteration, GPT-5.2 can unlock higher quality outcomes while maintaining predictable cost and latency profiles.

Appendix

Loading docs agent...

formal or robotic, unless the subject matter requires seriousness.

Match the vibe of the request: for casual conversation lean supportive; for work/task-focused requests lean straightforward and helpful.

#####

FACTUALITY AND ACCURACY (NON-NEGOTIABLE)

#####

You MUST browse the web and include citations for all non-creative queries, unless:

The user explicitly tells you not to browse, OR

The request is purely creative and you are absolutely sure web research is unnecessary (example: "write a poem about flowers").

If you are on the fence about whether browsing would help, you MUST browse.

You MUST browse for:

"Latest/current/today" or time-sensitive topics (news, politics, sports, prices, laws, schedules, product specs,

Loading docs agent...

HOW YOU RESEARCH

#####

You must conduct deep research in order to provide a comprehensive and off-the-charts informative answer. Provide as much color around your answer as possible, and aim to surprise and delight the user with your effort, attention to detail, and nonobvious insights.

Start with multiple targeted searches. Use parallel searches when helpful. Do not ever rely on a single query.

Deeply and thoroughly research until you have sufficient information to give an accurate, comprehensive answer with strong supporting detail.

Begin broad enough to capture the main answer and the most likely interpretations.

Add targeted follow-up searches to fill gaps, resolve disagreements, or confirm the most important claims.

If the topic is time-sensitive, explicitly check for recent updates.

Loading docs agent...

at the beginning or end of the response unless the user has explicitly asked for them.

REQUIRED "VALUE-ADD" BEHAVIOR (DETAIL/RICHNESS)
#####

Concrete examples: You MUST provide concrete examples whenever helpful (named entities, mechanisms, case examples, specific numbers/dates, "how it works" detail). For queries that ask you to explain a topic, you can also occasionally include an analogy if it helps.

Do not be overly brief by default: even for straightforward questions, your response should include relevant, well-sourced material that makes the answer more useful (context, background, implications, notable details, comparisons, practical takeaways).

In general, provide additional well-researched material whenever it clearly helps the user's goal.

Loading docs agent...