

Prompting is the process of providing input to a model. The quality of your output often depends on how well you're able to prompt the model.

Overview

Prompting is both an art and a science. OpenAI has some strategies and API design decisions to help you construct strong prompts and get consistently good results from a model. We encourage you to experiment.

Prompting tools and techniques

- [Prompt caching](#): Reuse stable prompt prefixes to reduce latency and input token costs on cache hits
- [Prompt engineering](#): Learn strategies, techniques, and tools to construct prompts

Refine your prompt

- Put overall tone or role guidance in the system message; keep task-specific details and examples in user messages.
- Combine few-shot examples into a concise YAML-style or bulleted block so your team can scan and update them.
- Mirror your project structure with clear folder names so teammates can locate prompts quickly.
- Run your prompt tests and evaluation cases every time you publish; catching issues early is cheaper than fixing them in production.

Prompts in your application

Treat prompts as application code. Store prompt content in named modules, build dynamic sections with typed function arguments, and review prompt changes in the same pull requests as the product behavior they support.



OpenAI is deprecating reusable prompt objects in the API. Prompt creation will be de-emphasized beginning June 3, 2026, and `v1/prompts` is scheduled to shut down on November 30, 2026. See the [deprecations page](#) for the current timeline.

Ask AI

For new work, don't create reusable prompt objects. Instead:

Loading docs agent...