

OpenAI offers two types of models: reasoning models (o3 and o4-mini, for example) and GPT models (like GPT-4.1). These model families behave differently.

This guide covers:

1. The difference between our reasoning and non-reasoning GPT models
2. When to use our reasoning models
3. How to prompt reasoning models effectively

Read more about reasoning models and how they work.

Reasoning models vs. GPT models

Compared to GPT models, our o-series models excel at different tasks and require different prompts. One model family isn't better than the other—they're just different.

We trained our o-series models ("the planners") to think longer and harder about complex tasks, making them effective at strategizing, planning solutions to complex problems, and making decisions based on large volumes of ambiguous information. These models can also execute tasks with high accuracy and precision, making them ideal for domains that would otherwise require a human expert—like math, science, engineering, financial services, and legal services.

On the other hand, our lower-latency, more cost-efficient GPT models ("the workhorses") are designed for straightforward execution. An application might use o-series models to plan out the strategy to solve a problem, and use GPT models to execute specific tasks, particularly when speed and cost are more important than perfect accuracy.

How to choose

What's most important for your use case?

- **Speed and cost** → GPT models are faster and tend to cost less
- **Executing well defined tasks** → GPT models handle explicitly defined tasks well
- **Accuracy and reliability** → o-series models are reliable decision makers
- **Complex problem-solving** → o-series models work through ambiguity and complexity

Ask AI

Loading docs agent...

When to use our reasoning models

Here are a few patterns of successful usage that we've observed from customers and internally at OpenAI. This isn't a comprehensive review of all possible use cases but, rather, some practical guidance for testing our o-series models.

[Ready to use a reasoning model? Skip to the quickstart →](#)

1. Navigating ambiguous tasks

Reasoning models are particularly good at taking limited information or disparate pieces of information and with a simple prompt, understanding the user's intent and handling any gaps in the instructions. In fact, reasoning models will often ask clarifying questions before making uneducated guesses or attempting to fill information gaps.

Loading docs agent...

We've found that reasoning models are particularly good at reasoning over complex documents that have hundreds of pages of dense, unstructured information—things like legal contracts, financial statements, and insurance claims. The models are particularly strong at drawing parallels between documents and making decisions based on unspoken truths represented in the data.

“Tax research requires synthesizing multiple documents to produce a final, cogent answer. We swapped GPT-4o for o1 and found that o1 was much better at reasoning over the interplay between documents to reach logical conclusions that were not evident in any one single document. As a result, we saw a 4x improvement in end-to-end performance by switching to o1—incredible.”

—[Blue J](#), AI platform for tax research”

Loading docs agent...

enabling other models to focus on execution.”


—[Argon AI](#), AI knowledge platform for the pharmaceutical industry”

““o1 powers many of our agentic workflows at Lindy, our AI assistant for work. The model uses function calling to pull information from your calendar or email and then can automatically help you schedule meetings, send emails, and manage other parts of your day-to-day tasks. We switched all of our agentic steps that used to cause issues to o1 and observing our agents becoming basically flawless overnight!”

—[Lindy AI](#), AI assistant for work”

5. Visual reasoning

Loading docs agent...



ABBREVIATIONS

#	POUNDS OR NUMBER	INT	INTERIOR
AC	AIR CONDITIONING	MAX	MAXIMUM
B.O.	BOTTOM OF	MIN	MINIMUM
B/T	BETWEEN	(N)	NEW
BD	BOARD	NTS	NOT TO SCALE
BM	BEAM	OC	ON CENTER
CLG	CEILING	PERF	PERFORATED
CLR	CLEAR	PT	PRESSURE-TREATED

2x8 Joists @ 24" O.C.	2x8	~500 linear ft	Wood	Floor framing for attic and porch
2x12 Joists @ 16" O.C.	2x12	~1,000 linear ft	Wood	Floor framing for main floors
4x4 PT Wood Post	4x4	~1	Pressure-Treated	Structural support in crawlspace and porch
4x12 Glulam Beam	4x12	~32 linear ft	Glulam	Main support beam in crawlspace
Anchor Bolts	1/2" diameter	~20	Steel	Used to secure sill plates to foundation

Loading docs agent...

7. Evaluation and benchmarking for other model responses

We've also seen reasoning models do well in benchmarking and evaluating other model responses. Data validation is important for ensuring dataset quality and reliability, especially in sensitive fields like healthcare. Traditional validation methods use predefined rules and patterns, but advanced models like o1 and o3-mini can understand context and reason about data for a more flexible and intelligent approach to validation.

“Many customers use LLM-as-a-judge as part of their eval process in Braintrust. For example, a healthcare company might summarize patient questions using a workhorse model like gpt-4o, then assess the summary quality with o1. One Braintrust customer saw the F1 score of a judge go from 0.12 with 4o to 0.74 with o1! In these use cases, they've found o1's reasoning to be a game-changer in

Loading docs agent...

examples first. If you have more complex requirements for your desired output, it may help to include a few examples of inputs and desired outputs in your prompt. Just ensure that the examples align very closely with your prompt instructions, as discrepancies between the two may produce poor results.

- **Provide specific guidelines:** If there are ways you explicitly want to constrain the model's response (like "propose a solution with a budget under \$500"), explicitly outline those constraints in the prompt.
- **Be very specific about your end goal:** In your instructions, try to give very specific parameters for a successful response, and encourage the model to keep reasoning and iterating until it matches your success criteria.
- **Markdown formatting:** Starting with `o1-2024-12-17`, reasoning models in the API will avoid generating responses with markdown formatting. To signal to the model when you do want markdown formatting in the response, include the string `Formatting re-enabled` on the first line of your developer message.

Loading docs agent...

performance regardless of the API being used.

Other resources

For more inspiration, visit the [OpenAI Cookbook](#), which contains example code and links to third-party resources, or learn more about our models and reasoning capabilities:

- [Meet the models](#)
- [Reasoning guide](#)
- [How to use reasoning for validation](#)
- [Video course: Reasoning with o1](#)

Loading docs agent...