

`gpt-live-1` is a voice model for natural, continuous conversation. It can listen and speak at the same time, respond to interruptions, and keep the conversation moving while a backend agent handles reasoning, tools, and longer tasks.

Use `session.instructions` for the assistant's role, speaking style, and when to ask the backend for help. Give the backend model or agent the procedures and tools for tasks such as looking up an order or changing a booking.

Describe the conversational behavior you want, and let GPT-Live choose the wording for ordinary replies. When migrating from Realtime, keep the rules your product needs for wording, interruptions, and the order of actions. Test the simpler prompt on representative conversations as you revise it.

Your application checks permissions and required confirmations before executing an action.

Recommended prompt structure

Start with this template and add instructions as needed. See [session configuration](#) for field limits.

Keep the template's `Backchannel policy`, `Interruption policy`, and `Delegation policy` headings, and customize the text beneath them. Backchannels are brief listening sounds, such as “mm-hmm,” that the assistant can make while the caller continues speaking.

You are [name], a calm, friendly voice assistant for [service].
Speak warmly and naturally, at an unhurried pace. Be clear and direct, not overly cheerful.
If the user is frustrated, acknowledge it briefly and focus on the next helpful step.

Backchannel policy: Use moderate backchannels. Acknowledge naturally without competing with the main response.

Interruption policy: Stop speaking when the user interrupts. Listen to what they say.

Delegation policy:

Backend tools:

– [capability]: [what the backend can do]



Ask AI

Loading docs agent...

Backchannel policy: Use moderate backchannels. Acknowledge naturally without competing with the main response.



If you want backchannels, allow brief listening sounds during interruptions. A rule that forbids all overlapping speech can suppress them.

Interruptions

When the user interrupts, the assistant should stop its answer and listen. A brief listening sound is different from taking over the user's turn.

Handle changes to a task separately from interruptions to speech. "Stop talking" asks the assistant to yield; "Cancel my booking" asks the backend to take an action. Have

Loading docs agent...

Delegate before giving an answer that depends on backend work.
Do not guess the result while waiting.

Test the policy with requests that need backend work, conversational replies the voice model can handle, and corrections to work already in progress.

Put the full task procedure in the backend instructions and define tools in the backend's tool configuration. Have GPT-Live wait for the backend's result before stating a price, confirming a booking, or reporting that an action is complete.

You can prompt GPT-Live to acknowledge a request while delegated work runs. As background work progresses, use `session.commentary.append` to provide updates you want GPT-Live to say aloud.

For backend prompts, conversation context, tool results, typed input, and API examples, read [Delegation and tools](#). For the architecture overview, read [Getting](#)

Loading docs agent...